

# 1. Pandas Series and DataFrame for Machine Learning

August 18, 2024

## 1 Introduction to the Pandas

Pandas is an open-source data manipulation and analysis library for Python, offering data structures and functions specifically designed to work with structured data seamlessly. Created by Wes McKinney in 2008, Pandas has become one of the most popular tools for data analysis in Python due to its ease of use, performance, and versatility.

At its core, Pandas provides two primary data structures:

1. **Series:** A one-dimensional array-like structure that can hold various data types (integers, floats, strings, etc.). Each element in a Series has an associated label, known as its index.
2. **DataFrame:** A two-dimensional table-like structure with labeled axes (rows and columns). It can be thought of as a collection of Series objects sharing the same index, allowing for a more complex data organization similar to a spreadsheet or SQL table.

### 1.1 Key Features of Pandas

- **Data Cleaning and Preparation:** Pandas makes it easy to handle missing data, filter unwanted entries, and transform data formats, ensuring that datasets are ready for analysis.
- **Data Alignment:** One of Pandas' most powerful features is automatic data alignment. When performing operations on data from different sources, Pandas automatically aligns the data based on the index, making it easier to combine and compare different datasets.
- **Efficient Data Manipulation:** Pandas provides a wide range of functions for data manipulation, such as merging, joining, reshaping, and pivoting datasets. These operations are optimized for performance, allowing you to work with large datasets efficiently.
- **Time Series Analysis:** Pandas has robust support for working with time series data, including date and time indexing, resampling, and rolling window calculations. This makes it an excellent tool for financial and economic data analysis.
- **Input/Output Tools:** Pandas can read and write data in various formats, including CSV, Excel, SQL databases, JSON, and more. This versatility makes it easy to integrate Pandas into existing data workflows.
- **Group By Operations:** Pandas allows you to split your data into groups based on specific criteria and then apply aggregate functions or transformations to each group independently.

## 1.2 Use Cases for Pandas

Pandas is widely used in various fields, including finance, economics, social sciences, engineering, and more. Some common use cases include:

- **Data Analysis:** Pandas is frequently used for exploring, cleaning, and analyzing datasets, whether for academic research, business intelligence, or machine learning.
- **Data Visualization:** While Pandas itself does not provide advanced plotting capabilities, it integrates well with libraries like Matplotlib and Seaborn, making it easier to visualize data trends and patterns.
- **Data Wrangling:** Pandas is often used for preparing and transforming data before feeding it into machine learning models or other statistical tools.

In summary, Pandas is a powerful and flexible tool that simplifies many aspects of data analysis and manipulation in Python. Whether you're working with small datasets or large, complex data, Pandas provides the tools you need to analyze, clean, and visualize your data efficiently.

## 1.3 Installation of Pandas

```
[ ]: !pip install pandas
```

```
Requirement already satisfied: pandas in  
c:\users\vishwajeet\appdata\local\programs\python\python312\lib\site-packages  
(2.2.2)  
Requirement already satisfied: numpy>=1.26.0 in  
c:\users\vishwajeet\appdata\local\programs\python\python312\lib\site-packages  
(from pandas) (2.0.1)  
Requirement already satisfied: python-dateutil>=2.8.2 in  
c:\users\vishwajeet\appdata\local\programs\python\python312\lib\site-packages  
(from pandas) (2.9.0.post0)  
Requirement already satisfied: pytz>=2020.1 in  
c:\users\vishwajeet\appdata\local\programs\python\python312\lib\site-packages  
(from pandas) (2024.1)  
Requirement already satisfied: tzdata>=2022.7 in  
c:\users\vishwajeet\appdata\local\programs\python\python312\lib\site-packages  
(from pandas) (2024.1)  
Requirement already satisfied: six>=1.5 in  
c:\users\vishwajeet\appdata\local\programs\python\python312\lib\site-packages  
(from python-dateutil>=2.8.2->pandas) (1.16.0)
```

## 1.4 Pandas Series Object

A **Pandas Series** is a one-dimensional labeled array capable of holding data of any type (integers, strings, floats, Python objects, etc.). It is similar to a column in a spreadsheet or a database table, but with more flexibility and functionality.

### 1.4.1 Key Characteristics of a Pandas Series

1. **One-Dimensional:**

- A Series is essentially a single column of data, making it one-dimensional. Unlike a list or a NumPy array, each element in a Series is associated with a label, known as an index.

## 2. Index:

- The index in a Series is a key feature that distinguishes it from other data structures like lists or arrays. Each element in a Series is indexed, meaning it has a label that allows you to access data based on a unique identifier rather than just its position. By default, Pandas assigns an integer index starting from 0, but you can customize the index to use labels, dates, or other identifiers.

## 3. Homogeneous Data:

- A Series can hold data of only one type, similar to an array in NumPy. However, it is more flexible because it can contain any data type, including numbers, strings, or even Python objects.

## 4. Data Alignment:

- One of the powerful features of a Series is automatic alignment based on the index labels. This means that when performing operations on two Series objects, Pandas automatically aligns them by their index, facilitating more intuitive data manipulation.

```
[ ]: # Importing all the libraries
import numpy as np
import pandas as pd
```

## 1.5 Creating a Series from Python Objects

```
[ ]: help(pd.Series)
```

## 1.6 Data Lists and Indexing

We can create a series from Python lists(also from NumPy arrays)

```
[ ]: city = ["Delhi", "Mumbai", "Kolkata", "Pune", "Bengaluru"]
```

```
[ ]: num = [10,20,30,40,50]
```

```
[ ]: series1 = pd.Series(data=city)
series1
```

```
[ ]: 0      Delhi
      1      Mumbai
      2      Kolkata
      3        Pune
      4    Bengaluru
dtype: object
```

```
[ ]: series2 = pd.Series(data=num)
series2
```

```
[ ]: 0    10
      1    20
```

```
2    30
3    40
4    50
dtype: int64
```

```
[ ]: series3 = pd.Series(data=num, index=city)
series3
```

```
[ ]: Delhi    10
Mumbai    20
Kolkata    30
Pune      40
Bengaluru  50
dtype: int64
```

Let's take another example

```
[ ]: age = np.random.randint(0,100,4)
age
```

```
[ ]: array([40, 60, 92, 64], dtype=int32)
```

```
[ ]: names = ["Raj", "Vishal", "Vinay", "Shekhar"]
names
```

```
[ ]: ['Raj', 'Vishal', 'Vinay', 'Shekhar']
```

```
[ ]: ages = pd.Series(age, names)
```

```
[ ]: ages
```

```
[ ]: Raj      40
Vishal     60
Vinay     92
Shekhar    64
dtype: int32
```

## 1.7 Pandas Series Object from Python Dictionary

```
[ ]: ages = {'Sammy':5, 'Frank':10, 'Spike':7}
```

```
[ ]: ages
```

```
[ ]: {'Sammy': 5, 'Frank': 10, 'Spike': 7}
```

```
[ ]: pd.Series(ages)
```

```
[ ]: Sammy      5
     Frank      10
     Spike       7
     dtype: int64
```

## 1.8 Named Index in Pandas Series

```
[ ]: # Imaginary Sales Data for 1st and 2nd Quarters for Global Company
q1 = {'Japan': 80, 'China': 450, 'India': 200, 'USA': 250}
q2 = {'Brazil': 100, 'China': 500, 'India': 210, 'USA': 260}
```

```
[ ]: # Convert into Pandas Series
sales_Q1 = pd.Series(q1)
sales_Q2 = pd.Series(q2)
```

```
[ ]: sales_Q1
```

```
[ ]: Japan      80
     China    450
     India     200
     USA       250
     dtype: int64
```

```
[ ]: # Call values based on Named Index
sales_Q1['Japan']
```

```
[ ]: np.int64(80)
```

```
[ ]: # Integer Based Location information also retained!
sales_Q1[0]
```

```
C:\Users\Vishwajeet\AppData\Local\Temp\ipykernel_14716\3172792608.py:2:
FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a
future version, integer keys will always be treated as labels (consistent with
DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
sales_Q1[0]
```

```
[ ]: np.int64(80)
```

### 1.8.1 Be careful with potential errors

```
[ ]: # Wrong Name
sales_Q1['France']
```

```
-----
KeyError                                Traceback (most recent call last)
File c:
  ↳ \Users\Vishwajeet\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\indexing.py:3805, in Index.get_loc(self, key)
  ↳ 
```

```

3804 try:
-> 3805     return self._engine.get_loc(casted_key)
3806 except KeyError as err:

File index.pyx:167, in pandas._libs.index.IndexEngine.get_loc()

File index.pyx:196, in pandas._libs.index.IndexEngine.get_loc()

File pandas\libs\hashtable_class_helper.pxi:7081, in pandas._libs.hashtable.
↳PyObjectHashTable.get_item()

File pandas\libs\hashtable_class_helper.pxi:7089, in pandas._libs.hashtable.
↳PyObjectHashTable.get_item()

```

**KeyError:** 'France'

The above exception was the direct cause of the following exception:

```

KeyError                                Traceback (most recent call last)
Cell In[112], line 2
      1 # Wrong Name
----> 2 sales_Q1['France']

```

```

File c:
↳\Users\Vishwajeet\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\series.py:1121, in Series.__getitem__(self, key)
    1118     return self._values[key]
    1120 elif key_is_scalar:
-> 1121     return self._get_value(key)
    1123 # Convert generator to list before going through hashable part
    1124 # (We will iterate through the generator there to check for slices)
    1125 if is_iterator(key):

```

```

File c:
↳\Users\Vishwajeet\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\series.py:1237, in Series._get_value(self, label, takeable)
    1234     return self._values[label]
    1236 # Similar to Index.get_value, but we do not fall back to positional
-> 1237 loc = self.index.get_loc(label)
    1239 if is_integer(loc):
    1240     return self._values[loc]

```

```

File c:
↳\Users\Vishwajeet\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\index.py:3812, in Index.get_loc(self, key)
    3807     if isinstance(casted_key, slice) or (
    3808         isinstance(casted_key, abc.Iterable)
    3809         and any(isinstance(x, slice) for x in casted_key)

```

```

3810     ):
3811         raise InvalidIndexError(key)
-> 3812     raise KeyError(key) from err
3813 except TypeError:
3814     # If we have a listlike key, _check_indexing_error will raise
3815     # InvalidIndexError. Otherwise we fall through and re-raise
3816     # the TypeError.
3817     self._check_indexing_error(key)

```

**KeyError:** 'France'

```

[ ]: # Accidental Extra Space
sales_Q1['USA ']

```

**KeyError** Traceback (most recent call last)

File c:

```

  ↳ \Users\Vishwajeet\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\indexing.py:3805, in Index.get_loc(self, key)
    3804 try:
-> 3805     return self._engine.get_loc(casted_key)
    3806 except KeyError as err:

```

File index.pyx:167, in pandas.\_libs.index.IndexEngine.get\_loc()

File index.pyx:196, in pandas.\_libs.index.IndexEngine.get\_loc()

File pandas\\_libs\hashtable\_class\_helper.pxi:7081, in pandas.\_libs.hashtable.PyObjectHashTable.get\_item()

File pandas\\_libs\hashtable\_class\_helper.pxi:7089, in pandas.\_libs.hashtable.PyObjectHashTable.get\_item()

**KeyError:** 'USA '

The above exception was the direct cause of the following exception:

**KeyError** Traceback (most recent call last)

Cell In[113], line 2

```

    1 # Accidental Extra Space
----> 2 sales_Q1['USA ']

```

File c:

```

  ↳ \Users\Vishwajeet\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\series.py:1121, in Series.__getitem__(self, key)
    1118     return self._values[key]
    1120 elif key_is_scalar:

```

```

-> 1121     return self._get_value(key)
    1123 # Convert generator to list before going through hashable part
    1124 # (We will iterate through the generator there to check for slices)
    1125 if is_iterator(key):

```

File c:

```

-> \Users\Vishwajeet\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\series.py:1237, in Series._get_value(self, label, takeable)
    1234     return self._values[label]
    1236 # Similar to Index.get_value, but we do not fall back to positional
-> 1237 loc = self.index.get_loc(label)
    1239 if is_integer(loc):
    1240     return self._values[loc]

```

File c:

```

-> \Users\Vishwajeet\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\index.py:3812, in Index.get_loc(self, key)
    3807     if isinstance(casted_key, slice) or (
    3808         isinstance(casted_key, abc.Iterable)
    3809         and any(isinstance(x, slice) for x in casted_key)
    3810     ):
    3811         raise InvalidIndexError(key)
-> 3812     raise KeyError(key) from err
    3813 except TypeError:
    3814     # If we have a listlike key, _check_indexing_error will raise
    3815     # InvalidIndexError. Otherwise we fall through and re-raise
    3816     # the TypeError.
    3817     self._check_indexing_error(key)

```

**KeyError:** 'USA '

```

[ ]: # Capitalization Mistake
sales_Q1['usa']

```

**KeyError** Traceback (most recent call last)

File c:

```

-> \Users\Vishwajeet\AppData\Local\Programs\Python\Python312\Lib\site-packages\pandas\core\index.py:3805, in Index.get_loc(self, key)
    3804 try:
-> 3805     return self._engine.get_loc(casted_key)
    3806 except KeyError as err:

```

File index.pyx:167, in pandas.\_libs.index.IndexEngine.get\_loc()

File index.pyx:196, in pandas.\_libs.index.IndexEngine.get\_loc()



```
File pandas\\_libs\\hashtable_class_helper.pxi:7081, in pandas._libs.hashtable.  
↳PyObjectHashTable.get_item()
```

```
File pandas\\_libs\\hashtable_class_helper.pxi:7089, in pandas._libs.hashtable.  
↳PyObjectHashTable.get_item()
```

**KeyError:** 'usa'

The above exception was the direct cause of the following exception:

**KeyError** Traceback (most recent call last)

Cell In[114], line 2

```
1 # Capitalization Mistake  
----> 2 sales_Q1['usa']
```

File c:

```
↳\\Users\\Vishwajeet\\AppData\\Local\\Programs\\Python\\Python312\\Lib\\site-packages\\pandas\\core\\series.py:1121, in Series.__getitem__(self, key)  
1118     return self._values[key]  
1120 elif key_is_scalar:  
-> 1121     return self._get_value(key)  
1123 # Convert generator to list before going through hashable part  
1124 # (We will iterate through the generator there to check for slices)  
1125 if is_iterator(key):
```

File c:

```
↳\\Users\\Vishwajeet\\AppData\\Local\\Programs\\Python\\Python312\\Lib\\site-packages\\pandas\\core\\series.py:1237, in Series._get_value(self, label, takeable)  
1234     return self._values[label]  
1236 # Similar to Index.get_value, but we do not fall back to positional  
-> 1237 loc = self.index.get_loc(label)  
1239 if is_integer(loc):  
1240     return self._values[loc]
```

File c:

```
↳\\Users\\Vishwajeet\\AppData\\Local\\Programs\\Python\\Python312\\Lib\\site-packages\\pandas\\core\\index.py:3812, in Index.get_loc(self, key)  
3807     if isinstance(casted_key, slice) or (  
3808         isinstance(casted_key, abc.Iterable)  
3809         and any(isinstance(x, slice) for x in casted_key)  
3810     ):  
3811         raise InvalidIndexError(key)  
-> 3812     raise KeyError(key) from err  
3813 except TypeError:  
3814     # If we have a listlike key, _check_indexing_error will raise  
3815     # InvalidIndexError. Otherwise we fall through and re-raise  
3816     # the TypeError.  
3817     self._check_indexing_error(key)
```

```
KeyError: 'usa'
```

## 1.9 Basic Operations with Pandas Series

```
[ ]: # Grab just the index keys
sales_Q1.keys()
```

```
[ ]: Index(['Japan', 'China', 'India', 'USA'], dtype='object')
```

```
[ ]: # Can Perform Operations Broadcasted across entire Series
sales_Q1 * 2
```

```
[ ]: Japan    160
      China    900
      India    400
      USA      500
      dtype: int64
```

```
[ ]: sales_Q2 / 100
```

```
[ ]: Brazil    1.0
      China     5.0
      India     2.1
      USA       2.6
      dtype: float64
```

## 1.10 Operations Between Two Series

```
[ ]: # Notice how Pandas informs you of mismatch with NaN
sales_Q1 + sales_Q2
```

```
[ ]: Brazil    NaN
      China    950.0
      India    410.0
      Japan    NaN
      USA      510.0
      dtype: float64
```

```
[ ]: # You can fill these with any value you want
sales_Q1.add(sales_Q2,fill_value=0)
```

```
[ ]: Brazil    100.0
      China    950.0
      India    410.0
      Japan     80.0
      USA      510.0
```

dtype: float64

### 1.10.1 Pandas DataFrame

A **Pandas DataFrame** is a two-dimensional, size-mutable, and potentially heterogeneous tabular data structure with labeled axes (rows and columns). It is one of the most widely used data structures in Pandas, providing a convenient way to store and manipulate data similar to a table in a relational database or an Excel spreadsheet.

#### Key Characteristics of a DataFrame

1. **Two-Dimensional:**

- A DataFrame is a two-dimensional structure, meaning it has rows and columns. Each column in a DataFrame is a Pandas Series, allowing you to work with data in a structured and organized manner.

2. **Labeled Axes:**

- The rows and columns in a DataFrame are labeled with an index (for rows) and column names (for columns). This labeling allows for intuitive data selection, filtering, and manipulation based on labels rather than just integer positions.

3. **Heterogeneous Data:**

- Unlike a Series, which holds data of a single type, a DataFrame can hold multiple data types across different columns. For example, one column could store integers, another could store strings, and yet another could store floating-point numbers.

4. **Size-Mutable:**

- DataFrames can grow or shrink as needed, meaning you can add or remove rows and columns without having to redefine the entire structure.

5. **Alignment:**

- Like Series, DataFrames automatically align data based on the labels when performing operations between different DataFrames or between a DataFrame and a Series. This alignment is done on both rows and columns.

**Creating a DataFrame** You can create a DataFrame in several ways:

- **From a Dictionary of Lists:**

```
import pandas as pd

data = {
    'Name': ['Alice', 'Bob', 'Charlie'],
    'Age': [25, 30, 35],
    'City': ['New York', 'San Francisco', 'Los Angeles']
}
df = pd.DataFrame(data)
print(df)
```

- **From a List of Dictionaries:**

```
data = [
    {'Name': 'Alice', 'Age': 25, 'City': 'New York'},
    {'Name': 'Bob', 'Age': 30, 'City': 'San Francisco'},
```

```

        {'Name': 'Charlie', 'Age': 35, 'City': 'Los Angeles'}
    ]
    df = pd.DataFrame(data)
    print(df)

```

- **From a 2D NumPy Array:**

```

import numpy as np

data = np.array([[25, 'Alice', 'New York'],
                 [30, 'Bob', 'San Francisco'],
                 [35, 'Charlie', 'Los Angeles']])
df = pd.DataFrame(data, columns=['Age', 'Name', 'City'])
print(df)

```
- **From a CSV File:**

```

df = pd.read_csv('data.csv')
print(df)

```

**Accessing Data in a DataFrame** You can access the data in a DataFrame using various methods:

- **Accessing Columns:**

```

print(df['Name']) # Returns the 'Name' column as a Series

```
- **Accessing Rows by Label:**

```

print(df.loc[0]) # Returns the first row as a Series

```
- **Accessing Rows by Integer Location:**

```

print(df.iloc[0]) # Same as above, using integer location

```
- **Accessing a Subset of Rows and Columns:**

```

print(df.loc[0:1, ['Name', 'Age']]) # Returns first two rows and 'Name' and 'Age' columns

```

**Operations on a DataFrame** DataFrames support a wide range of operations, including arithmetic operations, aggregation functions, and more complex data manipulations:

- **Element-Wise Operations:**

```

df['Age'] = df['Age'] + 1 # Increase each age by 1
print(df)

```
- **Filtering Data:**

```

filtered_df = df[df['Age'] > 30] # Filter rows where 'Age' > 30
print(filtered_df)

```
- **Aggregation Functions:**

```

mean_age = df['Age'].mean() # Calculate the mean of the 'Age' column
print(mean_age)

```

- **Adding and Dropping Columns:**

```
df['Country'] = 'USA' # Add a new column 'Country'
df = df.drop('City', axis=1) # Drop the 'City' column
print(df)
```

**DataFrame Methods** Pandas DataFrames come with numerous built-in methods for common data tasks:

- **Sorting:**

```
df = df.sort_values(by='Age', ascending=False) # Sort by 'Age' in descending order
```

- **Grouping:**

```
grouped = df.groupby('City').mean() # Group by 'City' and calculate the mean for each group
```

- **Merging and Joining:**

```
df1 = pd.DataFrame({'key': ['A', 'B', 'C'], 'value': [1, 2, 3]})
df2 = pd.DataFrame({'key': ['B', 'C', 'D'], 'value': [4, 5, 6]})
merged_df = pd.merge(df1, df2, on='key', how='inner') # Merge on 'key' with an inner join
print(merged_df)
```

The Pandas DataFrame is an essential tool for anyone working with structured data in Python. Its versatility, ease of use, and powerful built-in functions make it an invaluable resource for tasks ranging from simple data exploration to complex data analysis. Whether you are dealing with small datasets or large, complex data, Pandas DataFrames provide the tools you need to manage, manipulate, and analyze your data efficiently.

## 1.11 Import Libraries

```
[ ]: import numpy as np
import pandas as pd
```

## 1.12 Creating a DataFrame from Python Objects

```
[ ]: help(pd.DataFrame)
```

```
[ ]: ## Let's Create our first Pandas DataFrame
# Make sure the seed is in the same cell as the random call
np.random.seed(101)
df1 = np.random.randint(0,101,(4,3))
```

```
[ ]: df1
```

```
[ ]: array([[95, 11, 81],
          [70, 63, 87],
          [75,  9, 77],
          [40,  4, 63]], dtype=int32)
```

```
[ ]: indx = ['CA', 'NY', 'AZ', 'TX']
```

```
[ ]: cols = ['Jan', 'Feb', 'Mar']
```

```
[ ]: df = pd.DataFrame(data=df1)
df
```

```
[ ]:      0   1   2
0  95  11  81
1  70  63  87
2  75   9  77
3  40   4  63
```

```
[ ]: df = pd.DataFrame(data=df1, index=indx)
df
```

```
[ ]:      0   1   2
CA  95  11  81
NY  70  63  87
AZ  75   9  77
TX  40   4  63
```

```
[ ]: df = pd.DataFrame(data=df1, index=indx, columns=cols)
df
```

```
[ ]:      Jan  Feb  Mar
CA   95   11   81
NY   70   63   87
AZ   75    9   77
TX   40    4   63
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 4 entries, CA to TX
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Jan     4 non-null        int32
1   Feb     4 non-null        int32
2   Mar     4 non-null        int32
dtypes: int32(3)
memory usage: 80.0+ bytes
```

## 2 Reading a .csv file for a DataFrame

### 2.0.1 What is a CSV File?

A **CSV file** (Comma-Separated Values file) is a plain text file that stores tabular data (numbers and text) in a simple, structured format. Each line in a CSV file corresponds to a row in a table, and the values in each row are separated by a comma or another delimiter, such as a semicolon or tab.

#### Key Features of a CSV File

1. **Plain Text Format:**

- CSV files are human-readable plain text files. This makes them easy to create, edit, and share using basic text editors or more advanced spreadsheet programs like Microsoft Excel or Google Sheets.

2. **Simple Structure:**

- CSV files are structured in a straightforward manner. Each line represents a row in the table, and within each row, commas (or another delimiter) separate the values that correspond to the columns.

3. **No Data Types:**

- Unlike more complex file formats (like Excel or SQL databases), CSV files do not store information about data types, formatting, or complex structures. All values are stored as plain text, and it is up to the software reading the file to interpret the data types (e.g., integers, floats, strings).

4. **Compatibility:**

- CSV files are widely supported across different platforms, programming languages, and software tools. This makes them a popular choice for data exchange between different systems.

5. **Lack of Metadata:**

- CSV files typically do not contain metadata, such as information about the data source, column data types, or formatting details. This simplicity is one of the reasons CSV is so widely used, but it also means that additional context or processing may be required when working with the data.

**Example of a CSV File** A CSV file might look like this:

```
Name, Age, City
Alice, 25, New York
Bob, 30, San Francisco
Charlie, 35, Los Angeles
```

In this example: - The first line contains the column headers: “Name,” “Age,” and “City.” - Each subsequent line represents a row in the table, with values corresponding to the columns.

**Working with CSV Files** **Opening a CSV File:** - You can open CSV files with text editors like Notepad, or spreadsheet programs like Excel or Google Sheets, which will display the data in a tabular format.

**Saving Data as CSV:** - Most spreadsheet applications allow you to save data as a CSV file, making it easy to export data for use in other programs.

**Reading CSV Files in Python:** - Python's `pandas` library is often used to read CSV files into a `DataFrame`, making it easy to analyze and manipulate the data:

```
import pandas as pd

df = pd.read_csv('data.csv')
print(df)
```

**Writing to a CSV File:** - You can also write data to a CSV file using `pandas`:

```
df.to_csv('output.csv', index=False)
```

### Common Uses of CSV Files

- **Data Exchange:** CSV is commonly used for exporting data from one system and importing it into another, especially in situations where different software tools are used.
- **Data Storage:** CSV files are often used to store simple datasets that do not require complex formatting or data types.
- **Data Processing:** Many data processing tasks, especially in scripting or programming, involve reading data from CSV files, processing it, and then saving the results back to a CSV file.

CSV files are a simple, effective way to store and exchange tabular data. Their ease of use, compatibility with a wide range of software, and human-readable format make them a popular choice for data storage and transfer, especially when working with data across different platforms and programming environments.

### Print your current directory file path with `pwd`

```
[ ]: %pwd
```

```
[ ]: 'e:\\Tutorials\\Python_for_ML\\2. Pandas'
```

### List the files in your current directory with `ls`

```
[ ]: %ls
```

Volume in drive E has no label.

Volume Serial Number is 3CDB-88F2

Directory of e:\Tutorials\Python\_for\_ML\2. Pandas

```
08/18/2024  12:38 PM    <DIR>          .
08/18/2024  12:34 PM    <DIR>          ..
08/18/2024  12:38 PM    <DIR>          Datasets
08/18/2024  02:05 PM                763,414 Pandas for Machine Learning.ipynb
                1 File(s)                763,414 bytes
                3 Dir(s)  168,625,639,424 bytes free
```



## 2.1 Let's read the data from .csv file

```
[ ]: df = pd.read_csv('Datasets/tips.csv')
```

```
[ ]: df
```

```
[ ]:
   total_bill  tip  sex smoker  day  time  size  price_per_person  \
0      16.99  1.01 Female    No  Sun  Dinner     2           8.49
1      10.34  1.66  Male    No  Sun  Dinner     3           3.45
2      21.01  3.50  Male    No  Sun  Dinner     3           7.00
3      23.68  3.31  Male    No  Sun  Dinner     2          11.84
4      24.59  3.61 Female    No  Sun  Dinner     4           6.15
..      ...    ...    ...    ...  ...  ...    ...           ...
239     29.03  5.92  Male    No  Sat  Dinner     3           9.68
240     27.18  2.00 Female   Yes  Sat  Dinner     2          13.59
241     22.67  2.00  Male   Yes  Sat  Dinner     2          11.34
242     17.82  1.75  Male    No  Sat  Dinner     2           8.91
243     18.78  3.00 Female    No  Thur Dinner     2           9.39
```

	Payer Name	CC Number	Payment ID
0	Christy Cunningham	3560325168603410	Sun2959
1	Douglas Tucker	4478071379779230	Sun4608
2	Travis Walters	6011812112971322	Sun4458
3	Nathaniel Harris	4676137647685994	Sun5260
4	Tonya Carter	4832732618637221	Sun2251
..	...	...	...
239	Michael Avila	5296068606052842	Sat2657
240	Monica Sanders	3506806155565404	Sat1766
241	Keith Wong	6011891618747196	Sat3880
242	Dennis Dixon	4375220550950	Sat17
243	Michelle Hardin	3511451626698139	Thur672

[244 rows x 11 columns]

About this DataSet (in case you are interested)

- Description
  - One waiter recorded information about each tip he received over a period of a few months working in one restaurant. He collected several variables:
- Format
  - A data frame with 244 rows and 7 variables
- Details
  - tip in dollars,
  - bill in dollars,
  - sex of the bill payer,
  - whether there were smokers in the party,
  - day of the week,
  - time of day,
  - size of the party.

In all he recorded 244 tips. The data was reported in a collection of case studies for business statistics (Bryant & Smith 1995).

- References
  - Bryant, P. G. and Smith, M (1995) Practical Data Analysis: Case Studies in Business Statistics. Homewood, IL: Richard D. Irwin Publishing:
- Note: We created some additional columns with Fake data, including Name, CC Number, and Payment ID.

## 2.2 Obtaining Basic Information About DataFrame

```
[ ]: df.columns
```

```
[ ]: Index(['total_bill', 'tip', 'sex', 'smoker', 'day', 'time', 'size',  
          'price_per_person', 'Payer Name', 'CC Number', 'Payment ID'],  
          dtype='object')
```

```
[ ]: df.index
```

```
[ ]: RangeIndex(start=0, stop=244, step=1)
```

```
[ ]: df.head(3)
```

```
[ ]:      total_bill  tip  sex smoker  day  time  size  price_per_person \  
0      16.99  1.01  Female    No  Sun  Dinner    2           8.49  
1      10.34  1.66   Male    No  Sun  Dinner    3           3.45  
2      21.01  3.50   Male    No  Sun  Dinner    3           7.00
```

	Payer Name	CC Number	Payment ID
0	Christy Cunningham	3560325168603410	Sun2959
1	Douglas Tucker	4478071379779230	Sun4608
2	Travis Walters	6011812112971322	Sun4458

```
[ ]: df.tail(3)
```

```
[ ]:      total_bill  tip  sex smoker  day  time  size  price_per_person \  
241      22.67  2.00   Male   Yes  Sat  Dinner    2          11.34  
242      17.82  1.75   Male    No  Sat  Dinner    2           8.91  
243      18.78  3.00  Female    No  Thur  Dinner    2           9.39
```

	Payer Name	CC Number	Payment ID
241	Keith Wong	6011891618747196	Sat3880
242	Dennis Dixon	4375220550950	Sat17
243	Michelle Hardin	3511451626698139	Thur672

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 244 entries, 0 to 243
```

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	total_bill	244 non-null	float64
1	tip	244 non-null	float64
2	sex	244 non-null	object
3	smoker	244 non-null	object
4	day	244 non-null	object
5	time	244 non-null	object
6	size	244 non-null	int64
7	price_per_person	244 non-null	float64
8	Payer Name	244 non-null	object
9	CC Number	244 non-null	int64
10	Payment ID	244 non-null	object

dtypes: float64(3), int64(2), object(6)

memory usage: 21.1+ KB

```
[ ]: ## length of the dataframe  
len(df)
```

```
[ ]: 244
```

```
[ ]: df.describe()
```

```
[ ]:      total_bill      tip      size  price_per_person      CC Number  
count  244.000000  244.000000  244.000000      244.000000  2.440000e+02  
mean    19.785943    2.998279    2.569672        7.888197  2.563496e+15  
std      8.902412    1.383638    0.951100        2.914234  2.369340e+15  
min      3.070000    1.000000    1.000000        2.880000  6.040679e+10  
25%     13.347500    2.000000    2.000000        5.800000  3.040731e+13  
50%     17.795000    2.900000    2.000000        7.255000  3.525318e+15  
75%     24.127500    3.562500    3.000000        9.390000  4.553675e+15  
max     50.810000   10.000000    6.000000       20.270000  6.596454e+15
```

```
[ ]: df.describe().transpose()
```

```
[ ]:      count      mean      std      min  \  
total_bill  244.0  1.978594e+01  8.902412e+00  3.070000e+00  
tip         244.0  2.998279e+00  1.383638e+00  1.000000e+00  
size        244.0  2.569672e+00  9.510998e-01  1.000000e+00  
price_per_person  244.0  7.888197e+00  2.914234e+00  2.880000e+00  
CC Number    244.0  2.563496e+15  2.369340e+15  6.040679e+10  
  
      25%      50%      75%      max  
total_bill  1.334750e+01  1.779500e+01  2.412750e+01  5.081000e+01  
tip         2.000000e+00  2.900000e+00  3.562500e+00  1.000000e+01  
size        2.000000e+00  2.000000e+00  3.000000e+00  6.000000e+00  
price_per_person  5.800000e+00  7.255000e+00  9.390000e+00  2.027000e+01
```

```
CC Number          3.040731e+13  3.525318e+15  4.553675e+15  6.596454e+15
```

## 2.3 Selection and Indexing

Let's learn how to retrieve information from a DataFrame.

### 2.3.1 COLUMNS

We will begin by learning how to extract information based on the columns

```
[ ]: df.head()
```

```
[ ]:
total_bill  tip    sex smoker  day    time  size  price_per_person  \
0      16.99  1.01  Female    No  Sun  Dinner    2           8.49
1      10.34  1.66   Male    No  Sun  Dinner    3           3.45
2      21.01  3.50   Male    No  Sun  Dinner    3           7.00
3      23.68  3.31   Male    No  Sun  Dinner    2          11.84
4      24.59  3.61  Female    No  Sun  Dinner    4           6.15

      Payer Name      CC Number  Payment ID
0  Christy Cunningham  3560325168603410    Sun2959
1    Douglas Tucker   4478071379779230    Sun4608
2    Travis Walters   6011812112971322    Sun4458
3  Nathaniel Harris   4676137647685994    Sun5260
4    Tonya Carter    4832732618637221    Sun2251
```

#### Grab a Single Column

```
[ ]: df['total_bill']
```

```
[ ]: 0      16.99
      1      10.34
      2      21.01
      3      23.68
      4      24.59
      ...
      239    29.03
      240    27.18
      241    22.67
      242    17.82
      243    18.78
      Name: total_bill, Length: 244, dtype: float64
```

```
[ ]: type(df['total_bill'])
```

```
[ ]: pandas.core.series.Series
```

#### Grab Multiple Columns

```
[ ]: # Note how its a python list of column names! Thus the double brackets.
df[['total_bill', 'tip']]
```

```
[ ]:      total_bill  tip
0         16.99  1.01
1         10.34  1.66
2         21.01  3.50
3         23.68  3.31
4         24.59  3.61
..          ...   ...
239        29.03  5.92
240        27.18  2.00
241        22.67  2.00
242        17.82  1.75
243        18.78  3.00
```

[244 rows x 2 columns]

### Create New Columns

```
[ ]: df['tip_percentage'] = 100* df['tip'] / df['total_bill']
```

```
[ ]: df.head()
```

```
[ ]:      total_bill  tip  sex smoker  day  time  size  price_per_person \
0         16.99  1.01 Female    No  Sun  Dinner    2           8.49
1         10.34  1.66  Male    No  Sun  Dinner    3           3.45
2         21.01  3.50  Male    No  Sun  Dinner    3           7.00
3         23.68  3.31  Male    No  Sun  Dinner    2          11.84
4         24.59  3.61 Female    No  Sun  Dinner    4           6.15
```

	Payer Name	CC Number	Payment ID	tip_percentage
0	Christy Cunningham	3560325168603410	Sun2959	5.944673
1	Douglas Tucker	4478071379779230	Sun4608	16.054159
2	Travis Walters	6011812112971322	Sun4458	16.658734
3	Nathaniel Harris	4676137647685994	Sun5260	13.978041
4	Tonya Carter	4832732618637221	Sun2251	14.680765

```
[ ]: df['price_per_person'] = df['total_bill'] / df['size']
```

```
[ ]: df.head()
```

```
[ ]:      total_bill  tip  sex smoker  day  time  size  price_per_person \
0         16.99  1.01 Female    No  Sun  Dinner    2      8.495000
1         10.34  1.66  Male    No  Sun  Dinner    3      3.446667
2         21.01  3.50  Male    No  Sun  Dinner    3      7.003333
3         23.68  3.31  Male    No  Sun  Dinner    2     11.840000
4         24.59  3.61 Female    No  Sun  Dinner    4      6.147500
```

	Payer Name	CC Number	Payment ID	tip_percentage
0	Christy Cunningham	3560325168603410	Sun2959	5.944673
1	Douglas Tucker	4478071379779230	Sun4608	16.054159
2	Travis Walters	6011812112971322	Sun4458	16.658734
3	Nathaniel Harris	4676137647685994	Sun5260	13.978041
4	Tonya Carter	4832732618637221	Sun2251	14.680765

```
[ ]: help(np.round)
```

Help on `_ArrayFunctionDispatcher` in module `numpy`:

`round(a, decimals=0, out=None)`

Evenly round to the given number of decimals.

Parameters

-----

`a` : array\_like

Input data.

`decimals` : int, optional

Number of decimal places to round to (default: 0). If `decimals` is negative, it specifies the number of positions to the left of the decimal point.

`out` : ndarray, optional

Alternative output array in which to place the result. It must have the same shape as the expected output, but the type of the output values will be cast if necessary. See :ref:`ufuncs-output-type` for more details.

Returns

-----

`rounded_array` : ndarray

An array of the same type as ``a``, containing the rounded values. Unless ``out`` was specified, a new array is created. A reference to the result is returned.

The real and imaginary parts of complex numbers are rounded separately. The result of rounding a float is a float.

See Also

-----

`ndarray.round` : equivalent method

`around` : an alias for this function

`ceil`, `fix`, `floor`, `rint`, `trunc`

Notes

-----

For values exactly halfway between rounded decimal values, NumPy rounds to the nearest even value. Thus 1.5 and 2.5 round to 2.0, -0.5 and 0.5 round to 0.0, etc.

`np.round` uses a fast but sometimes inexact algorithm to round floating-point datatypes. For positive `decimals` it is equivalent to `np.true_divide(np rint(a * 10**decimals), 10**decimals)`, which has error due to the inexact representation of decimal fractions in the IEEE floating point standard [1]\_ and errors introduced when scaling by powers of ten. For instance, note the extra "1" in the following:

```
>>> np.round(56294995342131.5, 3)
56294995342131.51
```

If your goal is to print such values with a fixed number of decimals, it is preferable to use numpy's float printing routines to limit the number of printed decimals:

```
>>> np.format_float_positional(56294995342131.5, precision=3)
'56294995342131.5'
```

The float printing routines use an accurate but much more computationally demanding algorithm to compute the number of digits after the decimal point.

Alternatively, Python's builtin `round` function uses a more accurate but slower algorithm for 64-bit floating point values:

```
>>> round(56294995342131.5, 3)
56294995342131.5
>>> np.round(16.055, 2), round(16.055, 2) # equals 16.054999999999997
(16.06, 16.05)
```

## References

-----

.. [1] "Lecture Notes on the Status of IEEE 754", William Kahan,  
<https://people.eecs.berkeley.edu/~wkahan/ieee754status/IEEE754.PDF>

## Examples

-----

```
>>> np.round([0.37, 1.64])
array([0., 2.])
>>> np.round([0.37, 1.64], decimals=1)
array([0.4, 1.6])
>>> np.round([.5, 1.5, 2.5, 3.5, 4.5]) # rounds to nearest even value
array([0., 2., 2., 4., 4.])
>>> np.round([1,2,3,11], decimals=1) # ndarray of ints is returned
```

```
array([ 1,  2,  3, 11])
>>> np.round([1,2,3,11], decimals=-1)
array([ 0,  0,  0, 10])
```

### Adjust Existing Columns

```
[ ]: # Because pandas is based on numpy, we get awesome capabilities with numpy's
      ↪ universal functions!
df['price_per_person'] = np.round(df['price_per_person'],2)
```

```
[ ]: df.head()
```

```
[ ]:  total_bill  tip    sex smoker  day    time  size  price_per_person  \
0      16.99  1.01  Female    No  Sun  Dinner    2         8.49
1      10.34  1.66   Male    No  Sun  Dinner    3         3.45
2      21.01  3.50   Male    No  Sun  Dinner    3         7.00
3      23.68  3.31   Male    No  Sun  Dinner    2        11.84
4      24.59  3.61  Female    No  Sun  Dinner    4         6.15
```

	Payer Name	CC Number	Payment ID	tip_percentage
0	Christy Cunningham	3560325168603410	Sun2959	5.944673
1	Douglas Tucker	4478071379779230	Sun4608	16.054159
2	Travis Walters	6011812112971322	Sun4458	16.658734
3	Nathaniel Harris	4676137647685994	Sun5260	13.978041
4	Tonya Carter	4832732618637221	Sun2251	14.680765

### Remove Columns

```
[ ]: df = df.drop("tip_percentage",axis=1)
```

```
[ ]: df.head()
```

```
[ ]:  total_bill  tip    sex smoker  day    time  size  price_per_person  \
0      16.99  1.01  Female    No  Sun  Dinner    2         8.49
1      10.34  1.66   Male    No  Sun  Dinner    3         3.45
2      21.01  3.50   Male    No  Sun  Dinner    3         7.00
3      23.68  3.31   Male    No  Sun  Dinner    2        11.84
4      24.59  3.61  Female    No  Sun  Dinner    4         6.15
```

	Payer Name	CC Number	Payment ID
0	Christy Cunningham	3560325168603410	Sun2959
1	Douglas Tucker	4478071379779230	Sun4608
2	Travis Walters	6011812112971322	Sun4458
3	Nathaniel Harris	4676137647685994	Sun5260
4	Tonya Carter	4832732618637221	Sun2251



### 3 Index Basics

Before going over the same retrieval tasks for rows, let's build some basic understanding of the pandas DataFrame Index.

```
[ ]: df.head()
```

```
[ ]:   total_bill  tip  sex smoker  day  time  size  price_per_person  \
0      16.99  1.01 Female    No  Sun  Dinner    2           8.49
1      10.34  1.66  Male    No  Sun  Dinner    3           3.45
2      21.01  3.50  Male    No  Sun  Dinner    3           7.00
3      23.68  3.31  Male    No  Sun  Dinner    2          11.84
4      24.59  3.61 Female    No  Sun  Dinner    4           6.15
```

	Payer Name	CC Number	Payment ID
0	Christy Cunningham	3560325168603410	Sun2959
1	Douglas Tucker	4478071379779230	Sun4608
2	Travis Walters	6011812112971322	Sun4458
3	Nathaniel Harris	4676137647685994	Sun5260
4	Tonya Carter	4832732618637221	Sun2251

```
[ ]: df.index
```

```
[ ]: RangeIndex(start=0, stop=244, step=1)
```

```
[ ]: df.set_index('Payment ID')
```

```
[ ]:   total_bill  tip  sex smoker  day  time  size  \
Payment ID
Sun2959      16.99  1.01 Female    No  Sun  Dinner    2
Sun4608      10.34  1.66  Male    No  Sun  Dinner    3
Sun4458      21.01  3.50  Male    No  Sun  Dinner    3
Sun5260      23.68  3.31  Male    No  Sun  Dinner    2
Sun2251      24.59  3.61 Female    No  Sun  Dinner    4
...         ...  ...  ...    ...  ...  ...
Sat2657      29.03  5.92  Male    No  Sat  Dinner    3
Sat1766      27.18  2.00 Female   Yes  Sat  Dinner    2
Sat3880      22.67  2.00  Male   Yes  Sat  Dinner    2
Sat17        17.82  1.75  Male    No  Sat  Dinner    2
Thur672      18.78  3.00 Female    No  Thur Dinner    2
```

	price_per_person	Payer Name	CC Number
Payment ID			
Sun2959	8.49	Christy Cunningham	3560325168603410
Sun4608	3.45	Douglas Tucker	4478071379779230
Sun4458	7.00	Travis Walters	6011812112971322
Sun5260	11.84	Nathaniel Harris	4676137647685994
Sun2251	6.15	Tonya Carter	4832732618637221

```

...
Sat2657          9.68      Michael Avila  5296068606052842
Sat1766         13.59      Monica Sanders 3506806155565404
Sat3880         11.34      Keith Wong   6011891618747196
Sat17           8.91      Dennis Dixon   4375220550950
Thur672         9.39      Michelle Hardin 3511451626698139

```

[244 rows x 10 columns]

```
[ ]: df.head()
```

```

[ ]:   total_bill  tip  sex smoker  day  time  size  price_per_person  \
0      16.99  1.01  Female    No  Sun  Dinner    2             8.49
1      10.34  1.66   Male    No  Sun  Dinner    3             3.45
2      21.01  3.50   Male    No  Sun  Dinner    3             7.00
3      23.68  3.31   Male    No  Sun  Dinner    2            11.84
4      24.59  3.61  Female    No  Sun  Dinner    4             6.15

```

```

      Payer Name      CC Number Payment ID
0  Christy Cunningham  3560325168603410   Sun2959
1    Douglas Tucker   4478071379779230   Sun4608
2    Travis Walters   6011812112971322   Sun4458
3  Nathaniel Harris   4676137647685994   Sun5260
4    Tonya Carter    4832732618637221   Sun2251

```

```
[ ]: df = df.set_index('Payment ID')
```

```
[ ]: df.head()
```

```

[ ]:   total_bill  tip  sex smoker  day  time  size  \
Payment ID
Sun2959      16.99  1.01  Female    No  Sun  Dinner    2
Sun4608      10.34  1.66   Male    No  Sun  Dinner    3
Sun4458      21.01  3.50   Male    No  Sun  Dinner    3
Sun5260      23.68  3.31   Male    No  Sun  Dinner    2
Sun2251      24.59  3.61  Female    No  Sun  Dinner    4

```

```

      price_per_person      Payer Name      CC Number
Payment ID
Sun2959             8.49  Christy Cunningham  3560325168603410
Sun4608             3.45    Douglas Tucker   4478071379779230
Sun4458             7.00    Travis Walters   6011812112971322
Sun5260            11.84  Nathaniel Harris   4676137647685994
Sun2251             6.15    Tonya Carter    4832732618637221

```

```

[ ]: ## With the help of reset_index() we can reset the index
df = df.reset_index()

```

```
[ ]: df.head()
```

```
[ ]: Payment ID total_bill tip sex smoker day time size \
0 Sun2959 16.99 1.01 Female No Sun Dinner 2
1 Sun4608 10.34 1.66 Male No Sun Dinner 3
2 Sun4458 21.01 3.50 Male No Sun Dinner 3
3 Sun5260 23.68 3.31 Male No Sun Dinner 2
4 Sun2251 24.59 3.61 Female No Sun Dinner 4

price_per_person Payer Name CC Number
0 8.49 Christy Cunningham 3560325168603410
1 3.45 Douglas Tucker 4478071379779230
2 7.00 Travis Walters 6011812112971322
3 11.84 Nathaniel Harris 4676137647685994
4 6.15 Tonya Carter 4832732618637221
```

### 3.0.1 ROWS

Let's now explore these same concepts but with Rows.

```
[ ]: df.head()
```

```
[ ]: Payment ID total_bill tip sex smoker day time size \
0 Sun2959 16.99 1.01 Female No Sun Dinner 2
1 Sun4608 10.34 1.66 Male No Sun Dinner 3
2 Sun4458 21.01 3.50 Male No Sun Dinner 3
3 Sun5260 23.68 3.31 Male No Sun Dinner 2
4 Sun2251 24.59 3.61 Female No Sun Dinner 4

price_per_person Payer Name CC Number
0 8.49 Christy Cunningham 3560325168603410
1 3.45 Douglas Tucker 4478071379779230
2 7.00 Travis Walters 6011812112971322
3 11.84 Nathaniel Harris 4676137647685994
4 6.15 Tonya Carter 4832732618637221
```

```
[ ]: df = df.set_index('Payment ID')
```

```
[ ]: df.head()
```

```
[ ]: total_bill tip sex smoker day time size \
Payment ID
Sun2959 16.99 1.01 Female No Sun Dinner 2
Sun4608 10.34 1.66 Male No Sun Dinner 3
Sun4458 21.01 3.50 Male No Sun Dinner 3
Sun5260 23.68 3.31 Male No Sun Dinner 2
Sun2251 24.59 3.61 Female No Sun Dinner 4
```

	price_per_person	Payer Name	CC Number
Payment ID			
Sun2959	8.49	Christy Cunningham	3560325168603410
Sun4608	3.45	Douglas Tucker	4478071379779230
Sun4458	7.00	Travis Walters	6011812112971322
Sun5260	11.84	Nathaniel Harris	4676137647685994
Sun2251	6.15	Tonya Carter	4832732618637221

### Grab a Single Row

```
[ ]: # Index Number
df.iloc[0]
```

```
[ ]: total_bill      16.99
      tip            1.01
      sex            Female
      smoker          No
      day            Sun
      time           Dinner
      size            2
      price_per_person  8.49
      Payer Name      Christy Cunningham
      CC Number       3560325168603410
      Name: Sun2959, dtype: object
```

```
[ ]: # Name Based
df.loc['Sun2959']
```

```
[ ]: total_bill      16.99
      tip            1.01
      sex            Female
      smoker          No
      day            Sun
      time           Dinner
      size            2
      price_per_person  8.49
      Payer Name      Christy Cunningham
      CC Number       3560325168603410
      Name: Sun2959, dtype: object
```

### Grab Multiple Rows

```
[ ]: df.iloc[0:4]
```

```
[ ]:
      total_bill  tip  sex smoker  day  time  size \
Payment ID
Sun2959      16.99  1.01  Female    No  Sun  Dinner    2
Sun4608      10.34  1.66   Male    No  Sun  Dinner    3
Sun4458      21.01  3.50   Male    No  Sun  Dinner    3
```

Sun5260	23.68	3.31	Male	No	Sun	Dinner	2
---------	-------	------	------	----	-----	--------	---

  

	price_per_person	Payer Name	CC Number
Payment ID			
Sun2959	8.49	Christy Cunningham	3560325168603410
Sun4608	3.45	Douglas Tucker	4478071379779230
Sun4458	7.00	Travis Walters	6011812112971322
Sun5260	11.84	Nathaniel Harris	4676137647685994

```
[ ]: df.loc[['Sun2959','Sun5260']]
```

```
[ ]:
```

	total_bill	tip	sex	smoker	day	time	size	\
Payment ID								
Sun2959	16.99	1.01	Female	No	Sun	Dinner	2	
Sun5260	23.68	3.31	Male	No	Sun	Dinner	2	

  

	price_per_person	Payer Name	CC Number
Payment ID			
Sun2959	8.49	Christy Cunningham	3560325168603410
Sun5260	11.84	Nathaniel Harris	4676137647685994

**Remove Row** Typically are datasets will be large enough that we won't remove rows like this since we won't know their row location for some specific condition, instead, we drop rows based on conditions such as missing data or column values. The next lecture will cover this in a lot more detail.

```
[ ]: df.head()
```

```
[ ]:
```

	total_bill	tip	sex	smoker	day	time	size	\
Payment ID								
Sun2959	16.99	1.01	Female	No	Sun	Dinner	2	
Sun4608	10.34	1.66	Male	No	Sun	Dinner	3	
Sun4458	21.01	3.50	Male	No	Sun	Dinner	3	
Sun5260	23.68	3.31	Male	No	Sun	Dinner	2	
Sun2251	24.59	3.61	Female	No	Sun	Dinner	4	

  

	price_per_person	Payer Name	CC Number
Payment ID			
Sun2959	8.49	Christy Cunningham	3560325168603410
Sun4608	3.45	Douglas Tucker	4478071379779230
Sun4458	7.00	Travis Walters	6011812112971322
Sun5260	11.84	Nathaniel Harris	4676137647685994
Sun2251	6.15	Tonya Carter	4832732618637221

```
[ ]: df.drop('Sun2959',axis=0).head()
```

```
[ ]:      total_bill  tip    sex smoker  day    time  size  \
Payment ID
Sun4608      10.34  1.66   Male    No  Sun  Dinner    3
Sun4458      21.01  3.50   Male    No  Sun  Dinner    3
Sun5260      23.68  3.31   Male    No  Sun  Dinner    2
Sun2251      24.59  3.61  Female    No  Sun  Dinner    4
Sun9679      25.29  4.71   Male    No  Sun  Dinner    4
```

```
      price_per_person      Payer Name      CC Number
Payment ID
Sun4608           3.45    Douglas Tucker  4478071379779230
Sun4458           7.00    Travis Walters  6011812112971322
Sun5260          11.84  Nathaniel Harris  4676137647685994
Sun2251           6.15      Tonya Carter  4832732618637221
Sun9679           6.32      Erik Smith   213140353657882
```

**Insert a New Row** Pretty rare to add a single row like this. Usually you use `pd.concat()` to add many rows at once. You could use the `.append()` method with a list of `pd.Series()` objects, but you won't see us do this with realistic real-world data.

```
[ ]: one_row = df.iloc[0]
```

```
[ ]: one_row
```

```
[ ]: total_bill      16.99
tip                1.01
sex                Female
smoker             No
day                Sun
time              Dinner
size                2
price_per_person    8.49
Payer Name          Christy Cunningham
CC Number           3560325168603410
Name: Sun2959, dtype: object
```

```
[ ]: type(one_row)
```

```
[ ]: pandas.core.series.Series
```

```
[ ]: df.tail()
```

```
[ ]:      total_bill  tip    sex smoker  day    time  size  \
Payment ID
Sat2657      29.03  5.92   Male    No  Sat  Dinner    3
Sat1766      27.18  2.00  Female   Yes  Sat  Dinner    2
Sat3880      22.67  2.00   Male   Yes  Sat  Dinner    2
Sat17        17.82  1.75   Male    No  Sat  Dinner    2
```

Thur672	18.78	3.00	Female	No	Thur	Dinner	2
---------	-------	------	--------	----	------	--------	---

Payment ID	price_per_person	Payer Name	CC Number
Sat2657	9.68	Michael Avila	5296068606052842
Sat1766	13.59	Monica Sanders	3506806155565404
Sat3880	11.34	Keith Wong	6011891618747196
Sat17	8.91	Dennis Dixon	4375220550950
Thur672	9.39	Michelle Hardin	3511451626698139

## 4 Happy Programming!!!

[ ]: